

CS61A NOTE15 SQL

SQL Basics

SQL is an example of a declarative programming language. Statements do not describe computations directly, but instead describe the desired result of some computation. It is the role of the query interpreter of the database system to plan and perform a computational process to produce such a result.

Creating Tables

You can create SQL tables either from scratch 从头 or from existing tables 现有.

The following statement creates a table by specifying column names and values without referencing another table. Each `SELECT` clause specifies the values for one row, and `UNION` is used to join rows together. The `AS` clauses give a name to each column; it need not be repeated in subsequent rows after the first.

| berkeley | stanford | year |
|----------|----------|------|
| 30 | 7 | 2002 |
| 28 | 16 | 2003 |
| 17 | 38 | 2014 |

```
CREATE TABLE big_game AS
  SELECT 30 AS berkeley, 7 AS stanford, 2002 AS year UNION
  SELECT 28,           16,           2003 UNION
  SELECT 17,           38,           2014;
```

纯文本

SQL operators

Expressions in the `SELECT`, `WHERE`, and `ORDER BY` clauses can contain one or more of the following operators:

- comparison operators: `=`, `>`, `<`, `<=`, `>=`, `<>` or `!=` ("not equal")
- boolean operators: `AND`, `OR`
- arithmetic operators: `+`, `-`, `*`, `/`
- concatenation operator: `||`

纯文本

Output the ratio of Berkeley's score to Stanford's score each year:

```
select berkeley * 1.0 / stanford from big_game;
```

Output the sum of scores in years where both teams scored over 10 points:

```
select berkeley + stanford from big_game where berkeley > 10 and stanford >
```

Output a table with a single column and single row containing the value "he

```
SELECT "hello" || " " || "world";
```

Selecting From Tables

More commonly, we will create new tables by selecting specific columns that we want from existing tables by using a `SELECT` statement as follows:

纯文本

```
SELECT [columns] FROM [tables] WHERE [condition] ORDER BY [columns] LIMIT [
```

```
SELECT name,salary2023-salary2022 FROM salaries ORDER BY salary2023-salary2
```

records

| name | division | title | salary | supervisor |
|-----------------|----------------|--------------------|--------|-----------------|
| Ben Bitdiddle | Computer | Wizard | 60000 | Oliver Warbucks |
| Alyssa P Hacker | Computer | Programmer | 40000 | Ben Bitdiddle |
| Cy D Fect | Computer | Programmer | 35000 | Ben Bitdiddle |
| Lem E Tweakit | Computer | Technician | 25000 | Ben Bitdiddle |
| Louis Reasoner | Computer | Programmer Trainee | 30000 | Alyssa P Hacker |
| Oliver Warbucks | Administration | Big Wheel | 150000 | Oliver Warbucks |
| Eben Scrooge | Accounting | Chief Accountant | 75000 | Oliver Warbucks |
| Lana Lambda | Administration | Executive Director | 610000 | Lana Lambda |

纯文本

```
> SELECT "Ben" AS first, "Bitdiddle" AS last ; #select不用写出新表
Ben|Bitdiddle

> SELECT "Ben" AS first, "Bitdiddle" AS last UNION ; #合并到一张表
> SELECT "Louis", "Reasoner"
Ben|Bitdiddle
Louis|Reasoner
```

纯文本

```
#SELECT specific values from an existing table using a FROM clause

#SELECT ..., ..., FROM records
> SELECT name, division FROM records ;
Alyssa P Hacker|Computer
...
Robert Cratchet|Accounting

#SELECT * FROM 从表中选择所有列
> SELECT * FROM records ;
Alyssa P Hacker|Computer|Programmer|40000|Ben Bitdiddle
...
Robert Cratchet|Accounting|Scrivener|18000|Eben Scrooge

#SELECT, filter out rows using a WHERE clause, sort the resulting rows with

#SELECT [columns] FROM [tables] WHERE [condition] ORDER BY [criteri]
```

#升asc 降desc

```
> SELECT * FROM records WHERE title = "Programmer";
```

Alyssa P Hacker|Computer|Programmer|40000|Ben Bitdiddle

Cy D Fect|Computer|Programmer|35000|Ben Bitdiddle

```
> SELECT name, salary FROM records WHERE division = "Accounting" ORDER BY salary;
```

Eben Scrooge|75000

Robert Cratchet|18000

disc Q3: Oliver Employees

Write a query that outputs the **names** of employees that **Oliver Warbucks** directly **supervises** 监督

纯文本

```
SELECT name FROM records WHERE supervisor='Oliver Warbucks';#别漏分号
```

discQ4

Write a query that outputs **all information** about employees that **supervise themselves**.

纯文本

```
SELECT * FROM records WHERE name=supervisor; #SELECT *对应全部信息
```

discQ5: Rich Employees

Write a query that outputs the **names** of all employees with salary **greater than 50,000** **in alphabetical order**.

纯文本

```
SELECT name FROM records WHERE salary > 50000 ORDER BY name;
```

Q6: Raises

| salaries | | |
|-----------------|------------|------------|
| name | salary2022 | salary2023 |
| Ben Bitdiddle | 60000 | 80000 |
| Alyssa P Hacker | 40000 | 80000 |
| Cy D Fect | 35000 | 74000 |
| Lem E Tweakit | 25000 | 28000 |
| Louis Reasoner | 30000 | 30000 |
| Oliver Warbucks | 150000 | 120000 |
| Eben Scrooge | 75000 | 76000 |
| Robert Cratchet | 18000 | 20000 |
| Lana Lambda | 610000 | 610000 |

Write a query that outputs the **names** of the **top 3** employees with the largest **salary raises from 2022 to 2023** along with their **corresponding salary raises**, **ordered from largest to smallest** raise.(增长,desc)

纯文本

```
SELECT name,salary2023-salary2022 FROM salaries ORDER BY salary2023-salary2022  
#LIMIT 3的用法—取前三个，ORDER BY（可为运算）ASC/DESC，记得加；
```

Joins 合并

To select data from multiple tables, we can use joins. There are many types of joins, but the only one we'll worry about is the inner join. To perform an inner join on two or more tables, simply list them all out in the **FROM** clause of a **SELECT** statement:

纯文本

```
SELECT [columns] FROM [table1], [table2], ... WHERE [condition] ORDER BY [c
```

When we join two or more tables, the default output is a cartesian product. For example, if we joined `big_game` with `coaches`, we'd get the following:

合并默认笛卡尔积

| berkeley | stanford | year | name | start | end |
|----------|----------|------|---------------|-------|------|
| 30 | 7 | 2002 | Jeff Tedford | 2002 | 2012 |
| 28 | 16 | 2003 | Sonny Dykes | 2013 | 2016 |
| 17 | 38 | 2014 | Justin Wilcox | 2017 | null |

| berkeley | stanford | year | name | start | end |
|----------|----------|------|---------------|-------|------|
| 30 | 7 | 2002 | Jeff Tedford | 2002 | 2012 |
| 30 | 7 | 2002 | Sonny Dykes | 2013 | 2016 |
| 30 | 7 | 2002 | Justin Wilcox | 2017 | null |
| 28 | 16 | 2003 | Jeff Tedford | 2002 | 2012 |
| 28 | 16 | 2003 | Sonny Dykes | 2013 | 2016 |
| 28 | 16 | 2003 | Justin Wilcox | 2017 | null |
| 17 | 38 | 2014 | Jeff Tedford | 2002 | 2012 |
| 17 | 38 | 2014 | Sonny Dykes | 2013 | 2016 |
| 17 | 38 | 2014 | Justin Wilcox | 2017 | null |

纯文本

```
SELECT * FROM big_game, coaches WHERE year >= start AND year <= end;
```

```
17|38|2014|Sonny Dykes|2013|2016
28|16|2003|Jeff Tedford|2002|2012
30|7|2002|Jeff Tedford|2002|2012
```

```
SELECT name, year FROM big_game, coaches
```

```
...> WHERE berkeley > stanford AND year >= start AND year <= end;
```

```
Jeff Tedford|2003
```

```
Jeff Tedford|2002
```

In the queries above, none of the column names are ambiguous. For example, it is clear that the `name` column comes from the `coaches` table because there isn't a column in the `big_game` table with that name. However, if a column name exists in more than one of the tables being joined, or if we join a table with itself, we must disambiguate the column names using *aliases* 别名.

For examples, let's find out what the score difference is for each team between a game in `big_game` and any previous games. Since each row in this table represents one game, in order to compare two games we must join `big_game` with itself:有一个自己套用自己的过程

纯文本

```
SELECT b.Berkeley - a.Berkeley, b.Stanford - a.Stanford, a.Year, b.Year FROM
-11|22|2003|2014
-13|21|2002|2014
-2|9|2002|2003
```

SQL Aggregation

Previously, we have been dealing with queries that process one row at a time 一整行查询. When we join, we make pairwise combinations of all of the rows 组合行. When we use `WHERE`, we filter out certain rows based on the condition 选择一些行. Alternatively, applying an aggregate function such as `MAX(column)` combines the values in multiple rows.

By default 默认, we combine the values of the *entire* table. For example, if we wanted to count the number of flights from our `flights` table, we could use:

What if we wanted to group together the values in similar rows and perform the aggregation operations within those groups? We use a `GROUP BY` clause.

Here's another example. For each unique departure, collect all the rows having the same departure airport into a group. Then, select the `price` column and apply the `MIN` aggregation to recover the price of the cheapest departure from that group. The end result is a table of departure airports and the cheapest departing flight.

如果我们想把相似行中的数值分组,并在这些组中进行聚合操作,会怎么样呢? 我们使用 `GROUP BY` 子句。

Just like how we can filter out rows with `WHERE`, we can also filter out groups with `HAVING`. Typically, a `HAVING` clause should use an aggregation function. Suppose we want to see all airports with at least two departures:就像我们可以用 `WHERE` 过滤掉行一样，我们也可以用 `HAVING` 过滤掉组。通常，`HAVING` 子句应该使用一个聚合函数。假设我们想查看所有至少有两个航班起飞的机场：

Note that the `COUNT(*)` aggregate just counts the number of rows in each group. Say we want to count the number of *distinct* airports instead. Then, we could use the following query:注意，`COUNT(*)`聚合只是计算每组中的行数。假设我们想计算独立机场的数量。那么，我们可以使用下面的查询：

纯文本

```
SELECT COUNT(*) from FLIGHTS;  
13
```

```
SELECT departure, MIN(price) FROM flights GROUP BY departure;  
group by在此列对一些行操作  
AUH|932  
LAS|50  
LAX|89  
SEA|32  
SFO|40  
SLC|42
```

```
SELECT departure FROM flights GROUP BY departure HAVING COUNT(*) >= 2;  
having类似where但不是列而是行操作了  
LAX  
SFO  
SLC
```

```
SELECT COUNT(DISTINCT departure) FROM flights;  
计算不同值用count和distinct  
6
```


Summary

```
SELECT [columns]
      FROM [tables]
      WHERE [condition]
      GROUP BY [column]
      HAVING [condition]
      ORDER BY [columns] [ASC/DESC]
      LIMIT [limit];
```

Final - Summer 2022 Q8 (c)

studios: name, studio

movies: title, budget, boxoffice

(c) (6.0 points)

Using the **movies** and **studios** tables from the last 2 problems, write a query that find the top 2 studios with regards to average box office.

The table should be sorted from highest average to lowest. You should only include studios that have 2 or more movies made. The output table should have 2 columns, one for the studio and one for the average box office.

Your solution should work on any tables with the same columns.

```
SELECT _____
      FROM movies, studios WHERE _____
      GROUP BY _____ HAVING _____
      ORDER BY _____ DESC _____;
```

Final - Summer 2022 Q8 (c)

(c) (6.0 points)

Using the **movies** and **studios** tables from the last 2 problems, write a query that find the top 2 studios with regards to average box office.

The table should be sorted from highest average to lowest. You should only include studios that have 2 or more movies made. The output table should have 2 columns, one for the studio and one for the average box office.

Your solution should work on any tables with the same columns.

```
SELECT studio, AVG(boxoffice)
      FROM movies, studios WHERE name = title
      GROUP BY studio HAVING COUNT(*) >= 2
      ORDER BY AVG(boxoffice) DESC LIMIT 2;
```

Final - Spring 2018 Q7 (b)

7. (10 points) Gotta Select 'Em All

For the questions below, assume that the following two SQL statements have been executed. The `pokedex` table describes the names of some Pokémon and their heights in inches. The `evolve` table describes how those Pokémon can evolve into the other Pokémon in the `pokedex`.

```
CREATE TABLE pokedex AS
SELECT "Eevee" AS name, 12 AS height UNION
SELECT "Jolteon"      , 31      UNION
SELECT "Leafeon"      , 39      UNION
SELECT "Bulbasaur"    , 28      UNION
SELECT "Ivysaur"      , 39      UNION
SELECT "Venasaur"     , 79      UNION
SELECT "Charmander"   , 24      UNION
SELECT "Charmeleon"   , 43      UNION
SELECT "Charizard"    , 67;

CREATE TABLE evolve AS
SELECT "Eevee" AS before, "Jolteon" AS after UNION
SELECT "Eevee"      , "Leafeon"      UNION
SELECT "Bulbasaur"   , "Ivysaur"      UNION
SELECT "Ivysaur"     , "Venasaur"     UNION
SELECT "Charmander"  , "Charmeleon"  UNION
SELECT "Charmeleon"  , "Charizard";
```

Final - Spring 2018 Q7 (b)

- (b) (6 pt) Write a `SELECT` statement that results in a table with one row for each Pokémon that can evolve. The table should have two columns: the first contains the name of the Pokémon that can evolve, and the second contains the maximum increase in height that it can attain by evolving. For example, Eevee can grow as much as 27 inches (when evolving to Leafeon), so the result should contain the row ("Eevee", 27). Your statement should behave correctly even if the rows in `evolve` and `pokedex` were different. The result should only consider ways of evolving that are described by a single row in the `evolve` table.

```
SELECT before, max(b.height - a.height)

FROM evolve, pokedex AS a, pokedex AS b

WHERE before=a.name AND after=b.name

GROUP BY before;
```

Final - Summer 2017 Q10 - (a):

10. (18 points) I've brought ice and fire together.

Consider the following schema that represents users, products, and sales in a database management system.

```
create table users(uid, uname, date_created);
create table products(pid, pname, description, rating, price);
create table sales(time, pid, uid);
```

- uid (user ID), pid (product ID), rating, price are numbers while all other columns are strings.
- The uid uniquely identifies one user because there may be users with the same uname and date_created.
- The pid uniquely identifies one product because there may be products with the same column values.
- The uid and pid in each row of sales references a uid in users and a pid in products.

Express the following queries in SQL using only features we've covered in this course.

Recall: Rows can be ordered in either **asc** ending (increasing) or **desc** ending (decreasing) order.

- (a) (2 pt) Select the uname and product rating of any one user who purchased a highest-rated product (a product such that there is no other product rated higher). If there is more than one such product, return any one product.

```
select u.uname, p.rating
  from users as u, products as p, sales as s
 where u.uid = s.uid and s.pid = p.pid
 order by p.rating desc limit 1;
```

SQL!!!!

table name:
movies

2014 Movies
movie

| | budget | bo | sf |
|--------------|-------------------|-----------------------|---------------|
| 2014 Movies | Budget (millions) | Box Office (millions) | SciFi? (IMDb) |
| Kingsman | 81 | 406 | False |
| Unfriended | 1 | 54 | False |
| Interstellar | 165 | 672 | True |

Question 1

Write a query that outputs the budget that appear more than once in the table of movies.

Question 2

Write a query that will output a list of sci-fi movies ordered by the return of investment ratio $[(\text{box office} - \text{budget})/\text{budget}]$ from greatest to least.

Solution:

```
select budget from movies group by budget having count(*) >1;
```

Solution:

```
select movie from movies where sf = "true" order by -(bo-budget)/budget;
```